

Software Certification – Coding, Code, and Coders

Klaus Havelund and Gerard J. Holzmann

Laboratory for Reliable Software (LaRS)

Jet Propulsion Laboratory, California Institute of Technology

4800 Oak Grove Drive, Pasadena, California, 91109-8099

firstname.lastname@jpl.nasa.gov

ABSTRACT

We describe a certification approach for software development that has been adopted at our organization. JPL develops robotic spacecraft for the exploration of the solar system. The flight software that controls these spacecraft is considered to be *mission critical*. We argue that the goal of a software certification process cannot be the development of “*perfect*” software, i.e., software that can be formally proven to be correct under all imaginable and unimaginable circumstances. More realistically, the goal is to guarantee a software development process that is conducted by knowledgeable engineers, who follow generally accepted procedures to control known risks, while meeting agreed upon standards of workmanship. We target three specific issues that must be addressed in such a certification procedure: the coding process, the code that is developed, and the skills of the coders. The coding process is driven by standards. The code is mechanically checked against the standards with the help of state-of-the-art static source code analyzers. The coders, finally, are certified in on-site training courses that include formal exams.

Categories and Subject Descriptors

D.2.0 [Software Engineering]: General – *standards*. K.5.2. [Governmental Issues]: – *Regulation*. K.7.3: [The Computing Profession]: *Certification, Licensing and Testing*.

General Terms

Design, Reliability, Standardization, Verification, Legal Aspects.

Keywords

Coding standards, code review, static source code analysis, logic model checking, unit testing, safety- and mission-critical software.

1. INTRODUCTION

John Rushby once described the dilemma of current approaches to software verification or certification as follows: “*Because we cannot demonstrate how well we’ve done, we’ll show how hard we’ve tried.*” [1]. The statement is apt. Few, if any, organizations feel confident enough about their software development processes that they are willing to give an absolute guarantee of its fitness for use or so much as the absence of preventable flaws in workmanship.

As customers, and generally as users that have to rely on the safety and reliability of sometimes critically important software applications (e.g., as used by banks, car makers, or in medical

devices), we are quite used to the opposite: we are routinely asked to sign disclaimers that hold the software makers invulnerable to flaws in workmanship and all possible damage that might be caused by it. This in itself is remarkable.

At some point, perhaps a few decades ago, we may have expected that standard market principles would solve this problem: customers could have been expected to reject products that are delivered without warranties of fitness. But this is not what happened.

The driving principles that determine how software applications are developed and marketed give a significant advantage to the vendor who delivers a new service first, and merely commits to slowly improve while the product is in use, based on customer feedback. The customers, in this way, become part of what otherwise would be the testers, except this group of testers pays the vendor, instead of the reverse. As unsettling as this might be from a philosophical point of view, it works quite well for the vast majority of software products sold today.

A clear exception holds for the category of *safety-critical* and *mission-critical* software applications (the distinction is whether a system failure may result in death/serious injury to people versus loss of mission). Most will agree that different rules must apply here, since for these types of applications it cannot be considered acceptable for a vendor to decline all responsibility for the potential damage caused in return for a mere commitment to fix any problems not caught in the software development process until *after* they have manifested themselves to end-users. If we now look more carefully at *which* different rules are applied in these cases, we are in for a surprise. In many cases there are no software certification requirements, and those requirements that do exist can only be described as modest. Organizations are often only asked to show “how hard they’ve tried” and not that certain standards of workmanship are met.

As part of our research and work, we have inevitably gained experience with the analysis of many software products that are considered critical. At JPL this naturally includes the analysis of the control software for interplanetary spacecraft, but we have also been involved in a broad range of other types of safety-critical applications, including the investigation of specific aspects of automotive software (e.g., in the context of a study of the potential for sudden unintended acceleration of Toyota vehicles in 2010), medical device software, software used in the shutdown systems of nuclear power plants, railway signaling protocol software, etc.

NASA’s shuttle software [2] is often mentioned as an example of how critical software systems can reach a high level of safety and reliability. This software indeed has an exemplary track record of having a low residual fault density rate over the approximately three decades of use. Like any other human design, it is, of course,

Copyright 2011 Association for Computing Machinery. ACM acknowledges that this contribution was authored or co-authored by an employee, contractor or affiliate of the U.S. Government. As such, the Government retains a nonexclusive, royalty-free right to publish or reproduce this article, or to allow others to do so, for Government purposes only.

EMSOFT’11, October 9–14, 2011, Taipei, Taiwan.

Copyright 2011 ACM 978-1-4503-0714-7/11/10...\$10.00.

not completely free from defects, nor can it be expected to be. One could well say that the first principle adopted in the design of any system that is meant to be reliable is the recognition that no single system component can be perfect: every part has breaking points, some known and some unknown. Reliability and safety, therefore should be treated as system properties, not component properties. To complete the argument, in almost all cases of interest the software is merely one component in a larger system that includes also hardware and human operators as essential elements.

The software used for commercial airplanes, similarly, has an enviable track record for reliability. Again, the track record is not for perfection, because in any sufficiently complex system there are always residual defects that are discovered only after a system is delivered and goes into operation. The goal for certified software, therefore, cannot be to put a process in place that guarantees correctness under all circumstances – the goal is to produce a safe and reliable system that is built by competent, well-trained developers, following a process that controls risks and meets the evolving standard of skilled workmanship. In one sentence here, then we touch on three separate targets for a software certification process: the coding process that is followed, the code, and the coders. The certification process that is followed in the aerospace industry (e.g., for software used to control commercial airplanes), targets primarily the coding process, e.g. with standards such as DO-178B [3]. There are no strict requirements here for the use of specific verification tools, or for the certification of software developers themselves. The target is only to secure that due diligence was used in the development process itself.

2. CERTIFICATION PROCESS

At JPL, in the development of the control software for interplanetary spacecraft, such as the Mars Exploration Rovers [4], we have adopted a different process. The intent of this process is to subject not just the coding process, but also the code, and the coders, and to some extent the software managers as well, to some form of certification.

2.1 The Coding Process

For flight code, JPL has adopted an Institutional Coding Standard [5], with which it requires compliance in all newly developed mission-critical software written at JPL. Most flight software, by a significant margin is traditionally written in the C programming language, and therefore the JPL coding standard for flight software targets this language. The coding standard deliberately captures only risk-related rules for which compliance can be verified mechanically. Other than most other coding standards, then, this standard has real teeth: except in rare cases, non-compliance is not an option for our flight software developers. Because all rules in the standard are specifically risk related (i.e., we can often point at a mission anomaly or mission loss that was caused by the violation of the underlying principle), approval for non-compliance is also rarely requested or granted. An example of a risk-related rule in this coding standard is the abolition of all dynamic memory use and of recursive code. Some of the motivation for the rules can be traced to the Power of Ten rules, described in [6]. JPL further imposes fairly strict requirements on the code review process some of which is detailed in [7]. The coding standard for C is included in brief format in Appendix A, organized according to importance of the rules.

JPL has also developed a coding standard for Java [8]. Java is at JPL mostly used for ground software. That is, software that executes on ground-based computers in mission operation centers, for example receiving telemetry from and sending commands to the rovers. Although this is not embedded software, its correct behavior is important for the correct behavior of the rovers. The Java coding standard is in a draft form and is not yet an institutional standard. It is more liberal in certain areas than the C coding standard. For example, it does allow dynamic memory allocation (use of the **new** construct). The coding standard for Java is included in brief format in Appendix B, organized according to subject categories. The rules represent headlines, which in the full standard are explained in more detail.

2.2 The Code

The code is rightfully subject to the strictest requirements. Flight code, e.g. for the MSL mission [4], is checked nightly for compliance with the JPL C coding standard [5], and subjected to rigorous analysis with four separate state-of-the-art static source code analysis tools [7] (at the time of writing this includes the commercial tools [9]: coverity, codesonar, and odasa (from semmle), and the research tool uno). The warnings generated by each of these tools is combined with the output of mission-specific checkers that secure compliance with naming conventions, coding style, etc. In addition, all warnings, if any (there should be none), from the standard C compiler, used in pedantic mode with all warnings enabled, are included in the results that are provided to the software developers as part of the standard ‘scrub’ interface [7]. The developers are required to close out all reports before a formal code review is initiated. In peer code reviews, an additional source of input is provided by designated peer code reviewers, and added to the ‘scrub’ results.

Furthermore, each programmer is responsible for writing unit tests for his/her modules. A compilation build of the entire system includes running all unit tests, which have to succeed for the build to succeed. Separately, key parts of the software design are also checked for correctness and compliance with higher level design requirements with the help of logic model checkers, such as Spin [10]. Training in the use of logic model checkers is tacitly provided via (optional) graduate-level courses taught by members of the JPL Laboratory for Reliable Software in the Computer Science Department at the California Institute of Technology. Approximately ten JPL employees outside the Laboratory for Reliable Software have so far taken and passed this course, and have become proficient in the use of logic model checkers for the analysis and verification of flight software.

The Java coding standard in its entirety has in collaboration with semmle been implemented in semmle’s static code analyzer, and is currently being tested on JPL internal projects, including the MSL telemetry and command ground software. The implementation is being refined (in collaboration with semmle), driven by the results obtained during these tests. The refinement consists of finding the right balance between a low number of false positives and a high number of true negatives. Too many false positives will discourage use of the standard.

2.3 The Coders

Starting in 2010, JPL adopted a new procedure for the certification of flight software developers. The procedure itself is still subject to some revision, but once fully operational no software developer will be able to touch flight code (develop, manage, or modify) without having successfully completed a JPL specific Flight Software Certification course. The course consists

of three modules, focusing on (a) computing science principles, (b) JPL software development standard processes, and (c) software risk and software vulnerabilities. Each module takes two full days of instruction, for a total of six days for all three modules combined. Each module ends with an exam that must be completed with a passing grade. At the time of writing, the first twenty software developers have successfully completed this course, and have received their certificates. Others have not passed and will have to take the course, and the exams, again. New classes are held several times a year, until all software developers have been certified. At that point, we will likely add refresher courses for those who are already certified, in addition to the basic certification course itself, to keep pace with continuing developments in this field. The certification course intends to certify that all developers of critical code are familiar with basic computing science theory, and standard algorithms, are intimately familiar with the risks inherent in the use of the programming languages that are typically used for flight code, and understand not just the letter but also the rationale for the coding standard that they are expected to follow. The certification course also introduces developers to the tools (e.g., static analyzers, code review tools, and unit test tools) that they will be using in flight software development.

Perhaps as an aside, JPL has also instituted an (as yet non-required) course for senior management. Senior management normally has deep experience with spacecraft and mission design, but less so with software design principles. To date, most of JPL's senior management has taken and completed this course. The course is repeated once a year, and is by invitation only.

3. REGULATORY PROCESS

As noted in the introduction, members of our team have been involved in a broad range of software analysis applications, targeting not only aerospace but also safety-critical software used in automobiles, medical devices, and in the shutdown systems of nuclear power plants. It is perhaps noteworthy that at present there do not appear to be any strict regulatory requirements on the development of these critical types of software applications, neither on the code or the coders, on the organization that employs the coders, or on the processes that are followed in the coding process.

In the automotive industry there is reasonable consensus on at least one set of coding guidelines: the one developed by the organization MiraLtd, and known as the MISRA-C Coding Guidelines [11]. Curiously, although many developing organizations have publically expressed support for these guidelines, there is no requirement (or verification) that they actually comply with them.

Compliance with any reasonable standard, e.g., [5,6,11,8], can make it significantly simpler to analyze code for potential anomalies, and to revise, and maintain it longer term. Much the same is true in the medical device industry, where the FDA does not require compliance with any specific coding standard or software development process, and goes no further than to

recommend the use of state-of-the-art static source code analyzers as part of software development process, without actually requiring evidence that this is done. Similarly, the Nuclear Regulatory Commission has issued no comparable requirements for any software used in the shutdown systems of future nuclear power plants, nor does it seem to have plans to do so, as a key part of the licensing process.

We believe that in each of these cases the lack of requirements on software development is an omission that should be corrected. Where not following generally accepted principles for safe software development could be regarded as a lack of workmanship on the part of the developer or developing organization, with the potential effect of contributing to preventable software failure, inadequate regulation for safety-critical software systems that we all rely on could well be regarded as a failure of the regulatory process.

4. ACKNOWLEDGMENT

The research described in this paper was carried out at the Jet Propulsion Laboratory, California Institute of Technology, under a contract with the National Aeronautics and Space Administration.

5. REFERENCES

- [1] J. Rushby. Verified Software Systems – the Certification Perspective, <http://www.csl.sri.com/users/rushby/slides/vsr-roadmap-cert-apr06.pdf>.
- [2] C. Fishman. They Write the Right Stuff, <http://www.fastcompany.com/magazine/06/writestuff.html>.
- [3] RTCA-DO-178B, Software Considerations in Airborne Systems and Equipment Certification, December 1992.
- [4] Mars Science Laboratory mission (MSL), http://www.nasa.gov/mission_pages/msl/index.html.
- [5] JPL Institutional Coding Standard for the C Programming Language, http://lars-lab.jpl.nasa.gov/JPL_Coding_Standard_C.pdf.
- [6] The “Power of Ten” Coding Rules, <http://spinroot.com/p10/>.
- [7] G.J. Holzmann. Scrub: a Tool for Code Reviews, Innovations in Systems and Software Engineering, Vol. 6, No. 4, 2010, pp. 311-318.
- [8] JPL Draft Coding Standard for the Java Programming Language, http://lars-lab.jpl.nasa.gov/JPL_Coding_Standard_Java.pdf.
- [9] Static Source Code Analysis Tools, <http://spinroot.com/static/>.
- [10] G.J. Holzmann. The Spin Model Checker – Primer and Reference Manual, Addison-Wesley, 2004.
- [11] MISRA-C:2004, Guidelines for the Use of the C Language in Critical Systems, Mira Ltd.

APPENDICES

A – JPL Coding Standard for C

1 Language Compliance	
1	Do not stray outside the language definition.
2	Compile with all warnings enabled; use static source code analyzers.
2 Predictable Execution	
3	Use verifiable loop bounds for all loops meant to be terminating.
4	Do not use direct or indirect recursion.
5	Do not use dynamic memory allocation after task initialization.
*6	Use IPC messages for task communication.
7	Do not use task delays for task synchronization.
*8	Explicitly transfer write-permission (ownership) for shared data objects.
9	Place restrictions on the use of semaphores and locks.
10	Use memory protection, safety margins, barrier patterns.
11	Do not use goto, setjmp or longjmp.
12	Do not use selective value assignments to elements of an enum list.
3 Defensive Coding	
13	Declare data objects at smallest possible level of scope.
14	Check the return value of non-void functions, or explicitly cast to (void).
15	Check the validity of values passed to functions.
16	Use static and dynamic assertions as sanity checks.
*17	Use U32, I16, etc instead of predefined C data types such as int, short, etc.
18	Make the order of evaluation in compound expressions explicit.
19	Do not use expressions with side effects.
4 Code Clarity	
20	Make only very limited use of the C pre-processor.
21	Do not define macros within a function or a block.
22	Do not undefine or redefine macros.
23	Place #else, #elif, and #endif in the same file as the matching #if or #ifdef.
*24	Place no more than one statement or declaration per line of text.
*25	Use short functions with a limited number of parameters.
*26	Use no more than two levels of indirection per declaration.
*27	Use no more than two levels of dereferencing per object reference.
*28	Do not hide dereference operations inside macros or typedefs.
*29	Do not use non-constant function pointers.
30	Do not cast function pointers into other types.
31	Do not place code or declarations before an #include directive.
5 – MISRA <i>shall</i> compliance	
73 rules	All MISRA <i>shall</i> rules not already covered at Levels 1-4.
6 – MISRA <i>should</i> compliance	
*16 rules	All MISRA <i>should</i> rules not already covered at Levels 1-4.

All rules are ‘shall’ rules (must be followed), except those marked with ‘*’ which are ‘should’ rules (justified deviations allowed).

B – JPL Coding Standard for Java

1 – Process	
1	Compile with checks turned on.
2	Apply static analysis.
3	Document public elements.
4	Write unit tests.
2 – Names	
5	Use the standard naming conventions.
6	Do not override field or class names.
3 - Packages, Classes and Interfaces	
7	Make imports explicit.
8	Do not have cyclic package dependencies.
9	Obey the contract for equals().
10	Define both equals() and hashCode().
11	Define equals when adding fields.
12	Define equals with parameter type Object.
13	Do not use finalizers.
14	Do not implement the Cloneable interface.
15	Do not call nonfinal methods in constructors.
16	Select composition over inheritance.
4 – Fields	
17	Make fields private.
18	Do not use static mutable fields.
19	Declare immutable fields final.
20	Initialize fields before use.
5 – Methods	
21	Use assertions.
22	Use annotations.
23	Restrict method overloading.
24	Do not assign to parameters.
25	Do not return null arrays or collections.
26	Do not call System.exit.
6 - Declarations and Statements	
27	Have one concept per line.
28	Use braces in control structures.
29	Do not have empty blocks.
30	Use breaks in switch statements.
31	End switch statements with default.
32	Terminate if-else-if with else.
7 - Expressions	
33	Restrict side effects in expressions.
34	Use named constants for non-trivial literals.
35	Make operator precedence explicit.
36	Do not use reference equality.
37	Use only short-circuit logic operators.
38	Do not use octal values.
39	Do not use floating point equality.
40	Use one result type in conditional expressions.
41	Do not use string concatenation operator in loops.
8 - Exceptions	
42	Do not drop exceptions.
43	Do not abruptly exit a finally block.
9 - Types	
44	Use generics.
45	Use interfaces as types when available.
46	Use primitive types.
47	Do not remove literals from collections.
48	Restrict numeric conversions.
10 - Concurrency	
49	Program against data races.
50	Program against deadlocks.
51	Do not rely on the scheduler for synchronization.
52	Wait and notify safely.
11 - Complexity	
53	Reduce code complexity.